Image Processing
INT3404 20
Week 7: Feature extraction
Edge, Line, Texture

Lecturer: Nguyen Thi Ngoc Diep, Ph.D.

Email: ngocdiep@vnu.edu.vn

Schedule

Tuần	Nội dung	Yêu cầu đối với sinh viên (ngoài việc đọc tài liệu tham khảo)
1	Giới thiệu môn học	Cài đặt môi trường: Python 3, OpenCV 3, Numpy, Jupyter Notebook
2	Ảnh số (Digital image) – Phép toán điểm (Point operations) Làm quen với OpenCV + Python Điều chỉnh độ tương phản (Contrast adjust)– Ghép ảnh (Combining images)	
3	Histogram - Histogram equalization	Làm bài tập 1: điều chỉnh gamma tìm contrast hợp lý
4	Phép lọc trong không gian điểm ảnh (linear processing filtering)	Thực hành ở nhà
5	Phép lọc trong không gian điểm ảnh (linear processing filtering) (cont.) Thực hành: Cách tìm filters	Thực hành ở nhà
6	Thực hành: Ứng dụng của histogram; Tìm ảnh mẫu (Template matching)	Bài tập mid-term
7	Trích rút đặc trưng của ảnh Cạnh (Edge) và đường (Line) và texture	Thực hành ở nhà
8	Các phép biến đổi hình thái (Morphological operations)	Làm bài tập 2: tìm barcode
9	Chuyển đổi không gian – Miền tần số – Phép lọc trên miền tần số Thông báo liên quan đồ án môn học	Đăng ký thực hiện đồ án môn học
10	Xử lý ảnh màu (Color digital image)	Làm bài tập 3: Chuyển đổi mô hình màu và thực hiện phân vùng
11	Các phép biến đổi hình học (Geometric transformations)	Thực hành ở nhà
12	Nhiễu – Mô hình nhiễu – Khôi phục ảnh (Noise and restoration)	Thực hành ở nhà
13	Nén ảnh (Compression)	Thực hành ở nhà
14	Hướng dẫn thực hiện đồ án môn học	Trình bày đồ án môn học
15	Hướng dẫn thực hiện đồ án môn học Tổng kết cuối kỳ	Trình bày đồ án môn học

Recall

Week 7: Feature extraction

Edge, Line, Texture

Edge detection

What is an edge?

- An edge = a significant local change in the image intensity

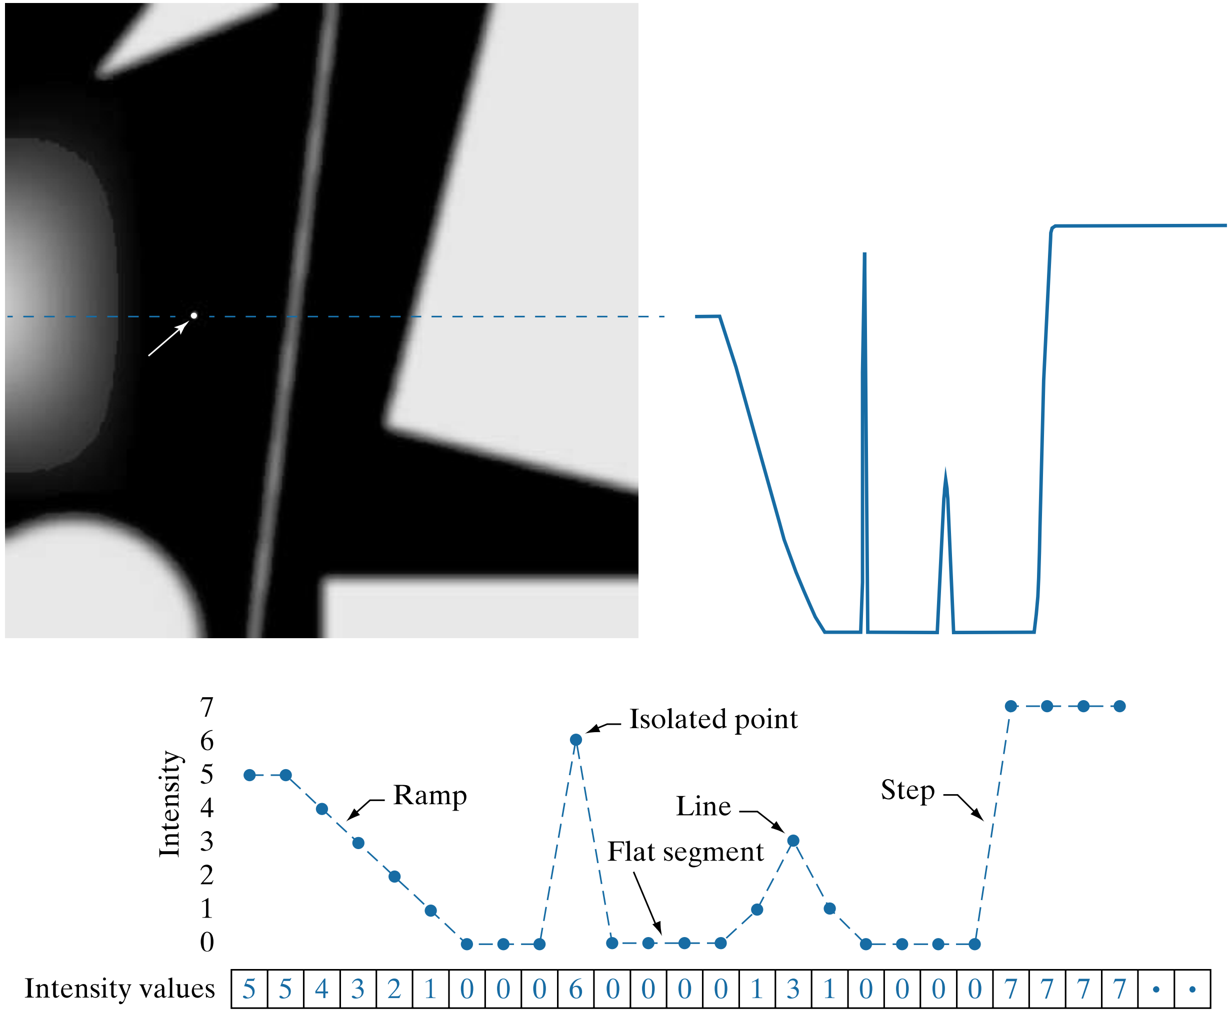

Image credit: Gonzalez et.al., Fig. 10.2

Image gradient

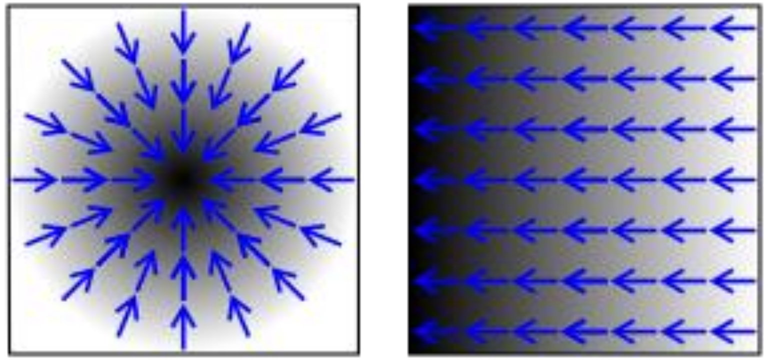

Two types of gradients, with blue arrows to indicate the direction of the gradient. Dark areas indicate higher values

Gradient

- Gradient of a function indicates how strong the function increases.
  - For 1-dimension function: $f(x) = x^2$

$$Grad(x) = \frac{\partial f(x)}{\partial (x)} = 2x$$

  - Grad(2)=4 indicates the the increasing direction of the function is to the right.
  - Grad(-1)=-2 indicates the increasing direction of the function is to the left.

First-order and second-order derivatives

$$\frac{\partial f}{\partial x} = f(x+1) - f(x)$$

$$\frac{\partial^2 f}{\partial x^2} = f(x+1) + f(x-1) - 2f(x)$$

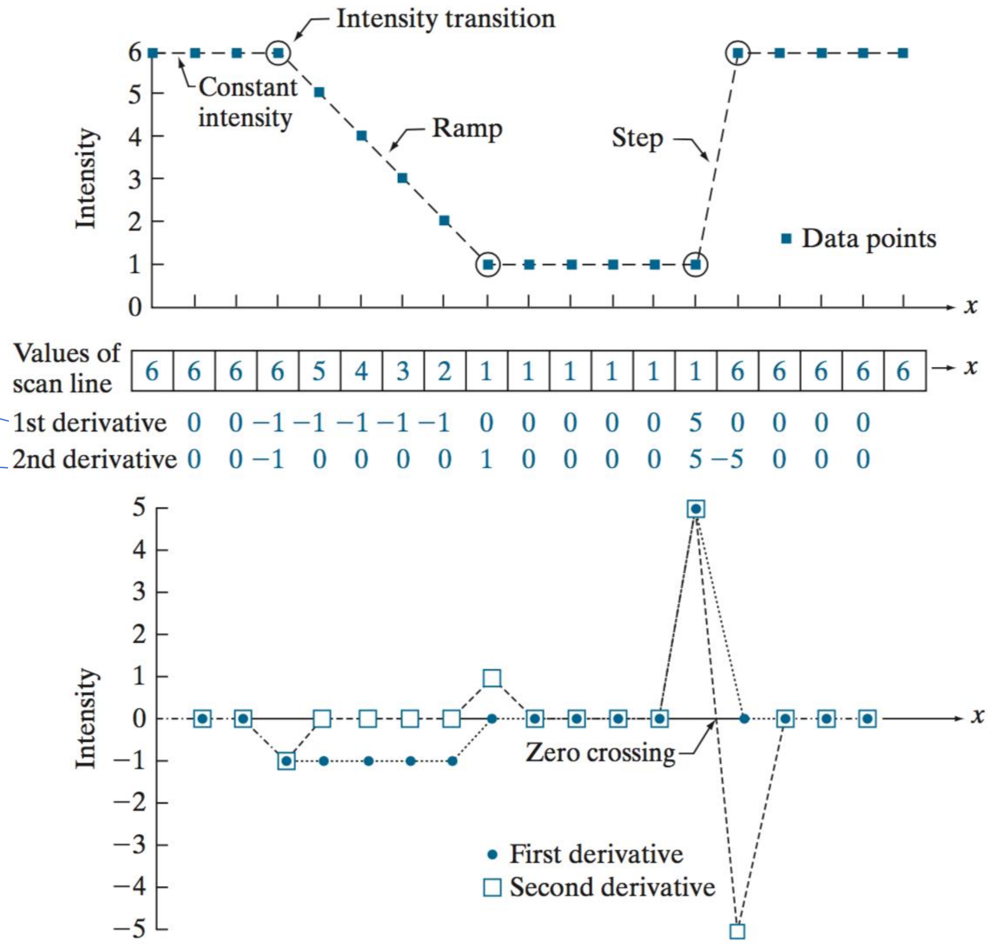

Image credit: Gonzalez, Fig. 3.44

Edge detection using derivatives

(1) Detecting the local maxima or minima of the first derivative

(2) Detecting the zero-crossings of the second derivative

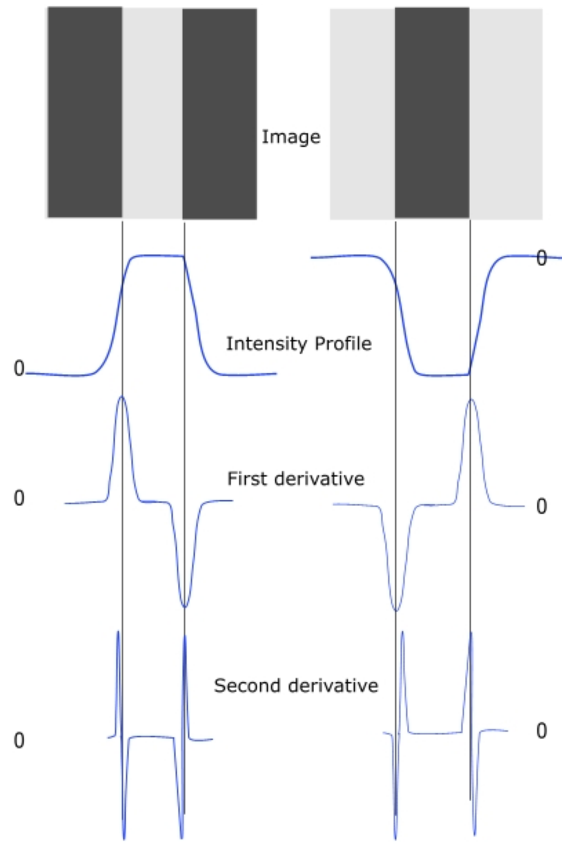

Image credit: MIPAV

Gradient of 2D discrete function

- Gradient of a 2-dimension function is calculated as follows:

$$Grad(x,y) = \frac{\partial f(x,y)}{\partial x}\vec{i} + \frac{\partial f(x,y)}{\partial y}\vec{j}$$

- The gradient is approximated as follows (first-order derivative) :

$$\frac{\partial f(x,y)}{\partial x} = f(x+1,y) - f(x,y), \frac{\partial f(x,y)}{\partial y} = f(x,y+1) - f(x,y)$$

Gradient of 2D image

$$\nabla f = \begin{bmatrix} g_x \\ g_y \end{bmatrix} = \begin{bmatrix} \frac{\partial f}{\partial x} \\ \frac{\partial f}{\partial y} \end{bmatrix}$$

convolution

$$\frac{\partial f}{\partial y} = \begin{bmatrix} -1 \\ +1 \end{bmatrix} * \mathbf{A}$$

Approximated by finite differences

$\Downarrow$

Avoid shifting of image

$$\begin{bmatrix} -1 \\ 0 \\ +1 \end{bmatrix}$$

Gradient

- The magnitude of gradient indicates the strong of edges:

$$\left|Grad(x,y)\right| = \sqrt{\left(\frac{\partial f(x,y)}{\partial y}\right)^2 + \left(\frac{\partial f(x,y)}{\partial x}\right)^2}$$

- Gradient computation procedure:
    - Calculate column gradient
    - Calculate row gradient
    - Calculate final gradient by the above function

Various kernels used to compute the gradient

z_1	z_2	z_3
z_4	z_5	z_6
z_7	z_8	z_9

A 3x3 region of an image

−1	0	0	−1
0	1	1	0

Roberts

$$g_x = \frac{\partial f}{\partial x} = (z_9 - z_5)$$
$$g_y = \frac{\partial f}{\partial y} = (z_8 - z_6)$$

−1	−1	−1	−1	0	1
0	0	0	−1	0	1
1	1	1	−1	0	1

Prewitt

$$g_x = \frac{\partial f}{\partial x} = (z_7 + z_8 + z_9) - (z_1 + z_2 + z_3)$$
$$g_y = \frac{\partial f}{\partial y} = (z_3 + z_6 + z_9) - (z_1 + z_4 + z_7)$$

−1	−2	−1	−1	0	1
0	0	0	−2	0	2
1	2	1	−1	0	1

Sobel

$$g_x = \frac{\partial f}{\partial x} = (z_7 + 2z_8 + z_9) - (z_1 + 2z_2 + z_3)$$
$$g_y = \frac{\partial f}{\partial y} = (z_3 + 2z_6 + z_9) - (z_1 + 2z_4 + z_7)$$

Pixel Difference masks

Column mask

$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & -1 \\ 0 & 0 & 0 \end{bmatrix}$$

Row mask

$$\begin{bmatrix} 0 & -1 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$$

Pixel difference example

Original Image

Column edges

Row edges

Final edges

Robert mask

- Roberts masks calculate gradient from two diagonals

Column

$$\begin{bmatrix} 0 & 0 & -1 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$$

Row

$$\begin{bmatrix} -1 & & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$$

Robert mask example

Original Image

Column edges

Row edges

Final edges

Prewitt mask

Column

$$\frac{1}{3}\begin{bmatrix} 1 & 0 & -1 \\ 1 & 0 & -1 \\ 1 & 0 & -1 \end{bmatrix}$$

Row

$$\frac{1}{3}\begin{bmatrix} -1 & -1 & -1 \\ 0 & 0 & 0 \\ 1 & 1 & 1 \end{bmatrix}$$

Prewitt mask

Original Image

Row edges

Column edges

Final edges

Sobel mask

Column

$$\frac{1}{4}\begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix}$$

Row

$$\frac{1}{4}\begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$

$I * f \overset{?}{=} I'$

0

value

= size

padding

default

Sobel filter example

Original Image

Column

Row

Combination

Laplace gradient

- Laplace edge in a continuous domain

$$G(x,y) = -\left(\frac{\partial^2 f(x,y)}{\partial x^2} + \frac{\partial^2 f(x,y)}{\partial y^2} \right)$$

- In a discrete domain, Laplace edge is approximated by

$$\begin{aligned} G(x,y) = &[f(x,y) - f(x,y-1] - [f(x,y+1) - f(x,y)] \\ &+ [f(x,y) - f(x+1,y] - [f(x-1,y) - f(x,y)] \end{aligned}$$

$$= f(x,y) * H(x,y)$$

Laplace mask

$$H = \begin{bmatrix} 0 & 0 & 0 \\ -1 & 2 & -1 \\ 0 & 0 & 0 \end{bmatrix} + \begin{bmatrix} 0 & -1 & 0 \\ 0 & 2 & 0 \\ 0 & -1 & 0 \end{bmatrix}$$

$$= \begin{bmatrix} 0 & -1 & 0 \\ -1 & 4 & -1 \\ 0 & -1 & 0 \end{bmatrix}$$

Laplace filter example

Original image I

|I*H|

Highboost filtering with Laplace

- Overall

$$I_{highboost} = c \cdot I_{orginal} + I_{highpass}$$

$$= \left(c \cdot \begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix} + H \right) * I_{original}$$

- Using Laplace mask

$$I_{highboost} = \left(c \cdot \begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix} + \begin{bmatrix} 0 & -1 & 0 \\ -1 & 4 & -1 \\ 0 & -1 & 0 \end{bmatrix} \right) * I_{original} = \begin{bmatrix} 0 & -1 & 0 \\ -1 & c+4 & -1 \\ 0 & -1 & 0 \end{bmatrix} * I_{original}$$

Highboost filter example

Original Image

c=0.5

c=1

Gradient comparison

Pixel difference

Robert

Prewitt

Sobel

Laplace

More advanced edge detection

LoG edge detection

$$\nabla^2 G(x,y) = \left(\frac{x^2 + y^2 - 2\sigma^2}{\sigma^4}\right) e^{-\frac{x^2+y^2}{2\sigma^2}}$$

$\sigma = 1.4$

$$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 4 & -1 \\ 0 & -1 & 0 \end{bmatrix}$$

$$\begin{pmatrix} 0 & 0 & 3 & 2 & 2 & 2 & 3 & 0 & 0 \\ 0 & 2 & 3 & 5 & 5 & 5 & 3 & 2 & 0 \\ 3 & 3 & 5 & 3 & 0 & 3 & 5 & 3 & 3 \\ 2 & 5 & 3 & -12 & -23 & -12 & 3 & 5 & 2 \\ 2 & 5 & 0 & -23 & -40 & -23 & 0 & 5 & 2 \\ 2 & 5 & 3 & -12 & -23 & -12 & 3 & 5 & 2 \\ 3 & 3 & 5 & 3 & 0 & 3 & 5 & 3 & 3 \\ 0 & 2 & 3 & 5 & 5 & 5 & 3 & 2 & 0 \\ 0 & 0 & 3 & 2 & 2 & 2 & 3 & 0 & 0 \end{pmatrix}$$

LoG edge detection example

Original image

LoG filtered

LoG edge

Canny edge detection

1. Smooth with 5x5 Gaussian kernel

2. Gradient with Sobel kernels

$$Edge\_Gradient\ (G) = \sqrt{G_x^2 + G_y^2}$$

$$Angle\ (\theta) = \tan^{-1}\left(\frac{G_y}{G_x}\right)$$

3. Non-maximum suppression

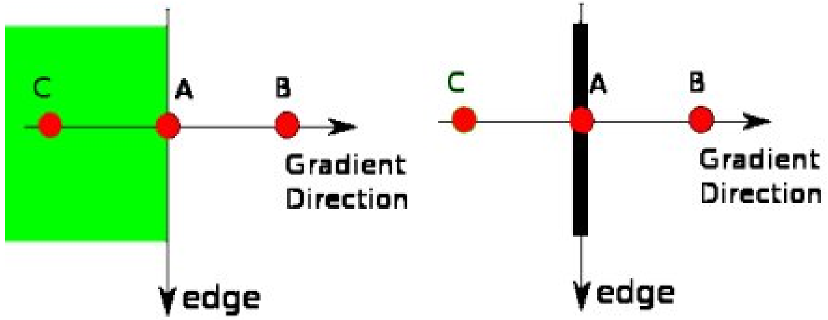

4. Thresholding

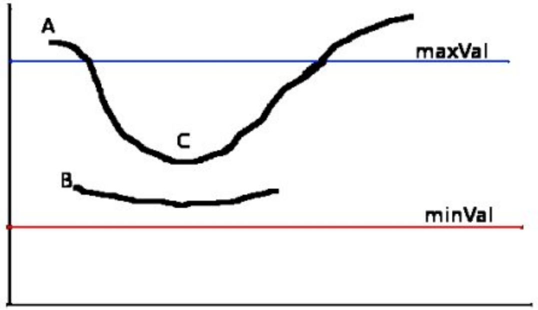

https://towardsdatascience.com/canny-edge-detection-step-by-step-in-python-computer-vision-b49c3a2d8123

Canny edge detection example

Original image

Canny

Line detection

Hough transform

Example: Line fitting

- Many objects characterized by presence of straight lines

line
thin edge

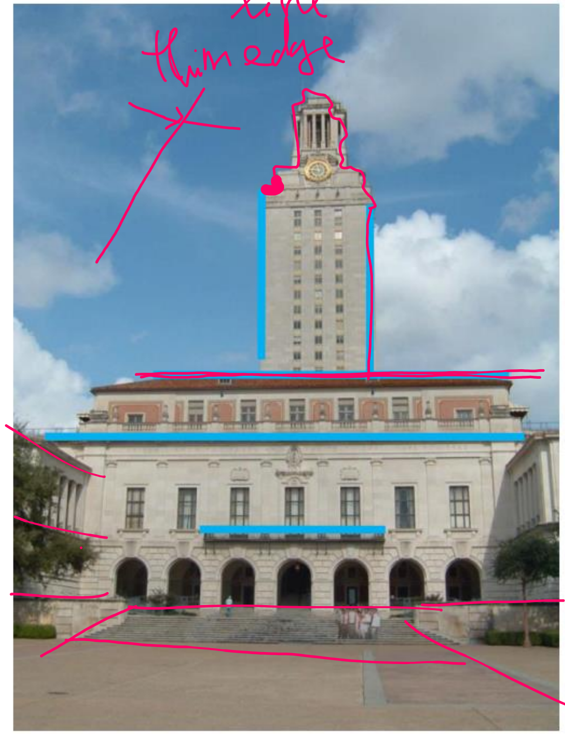

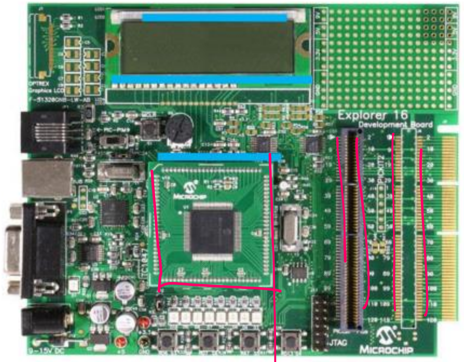

Difficulty of line fitting

edge pixel

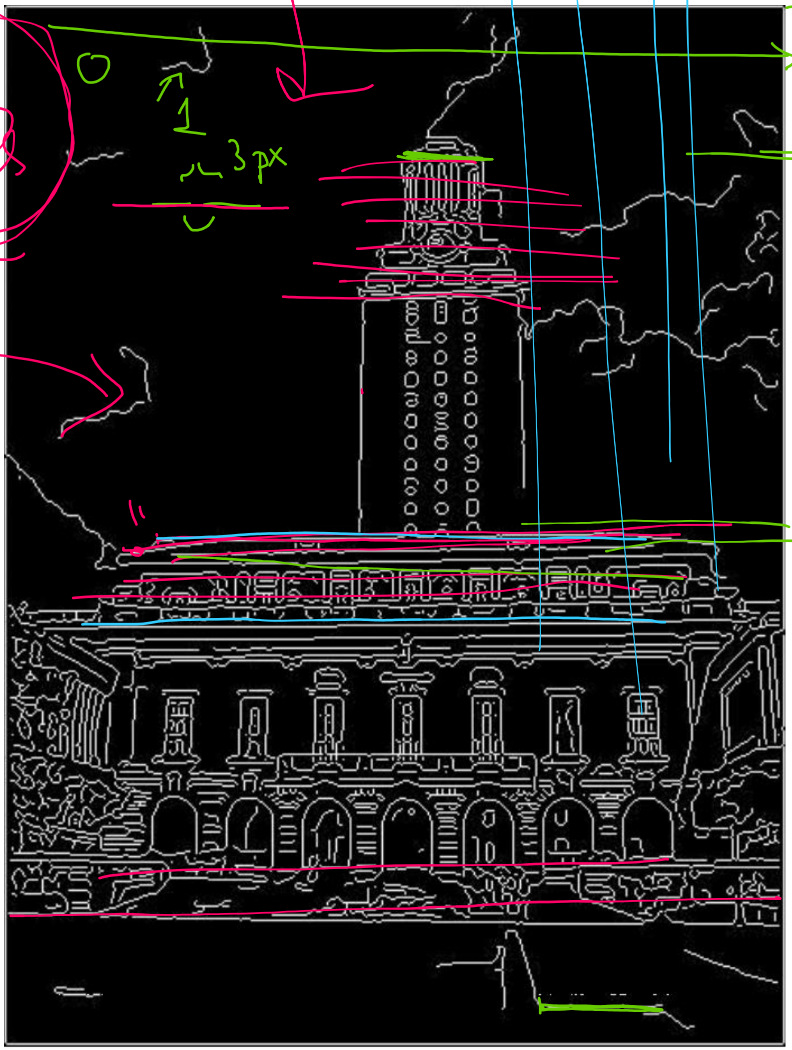

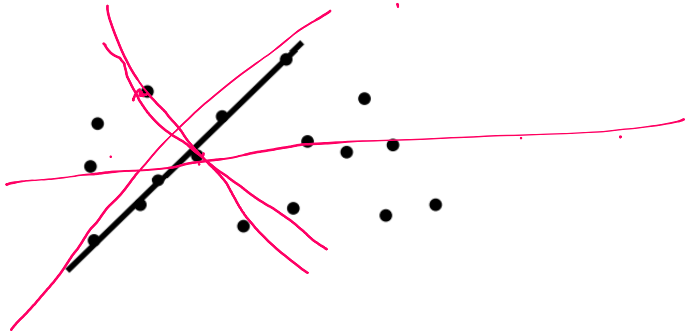

- Extra edge points (clutter), multiple models
  - Which points go with which line, if any?
- Only some parts of each line detected, and some parts are missing:
  - How to find a line that bridges missing evidence?
- Noise in measured edge points, orientations:
  - How to detect true underlying parameters

vertical / horizontal profile

Fitting lines with Hough transform

- Given points that belong to a line, what is the line?
- How many lines are there?
- Which points belong to which lines?
- Hough transform is a voting technique that can be used to answer all of these questions
- Main idea:
    1. Record vote for each possible line on which each edge point lies
    2. Look for lines that get many votes

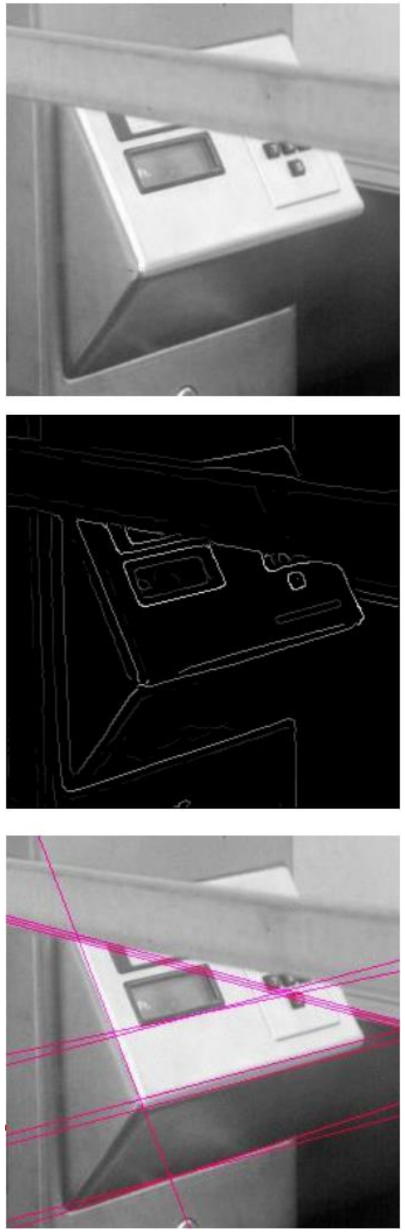

Line planes

$y = b$

ab-plane or parameter space

a b

FIGURE 10.28
(a) xy-plane.
(b) Parameter space.

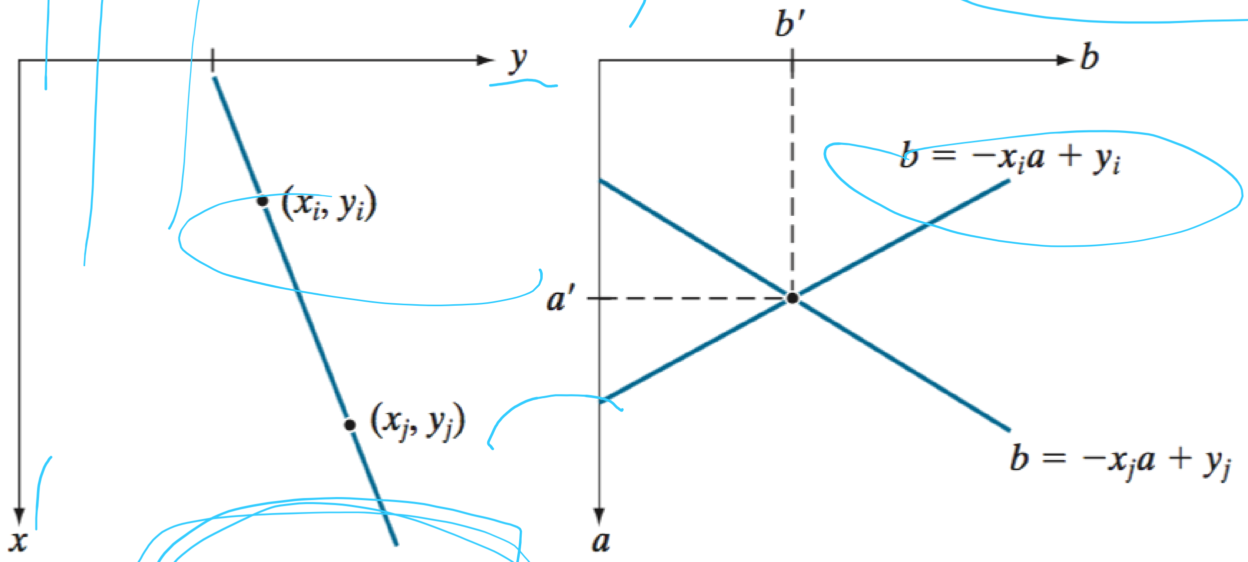

y = a'x + b'

What if the line approaches the vertical or horizontal direction?
(i.e., infinity slope)

Polar representation for lines

$$(\vec{P}-\vec{P_0})\cdot\vec{P_0} = 0$$

$$\vec{P}\cdot\vec{P_0} = \vec{P_0}\cdot\vec{P_0}$$

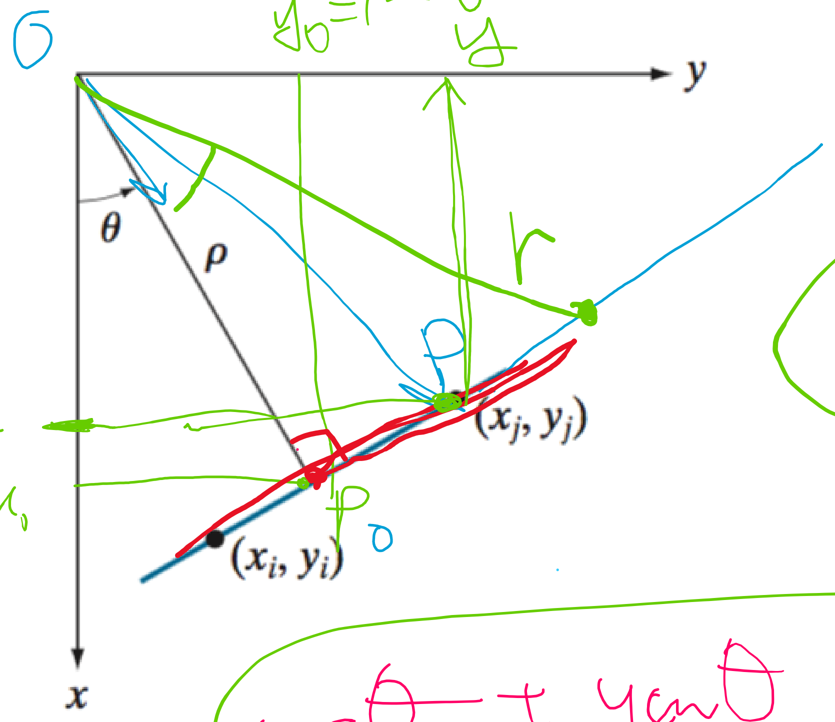

$$x\cos\theta + y\sin\theta = \rho$$

→ Point in image space is now sinusoid segment in Hough space

① intuition

Finding lines in an image: Hough algorithm

Polar voting

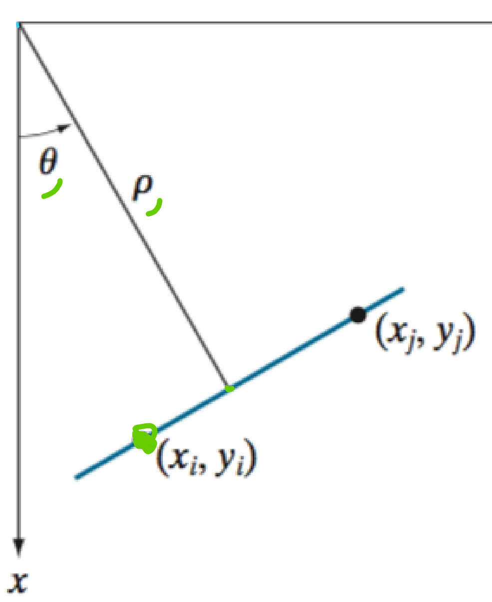

1. Using the polar parameterization

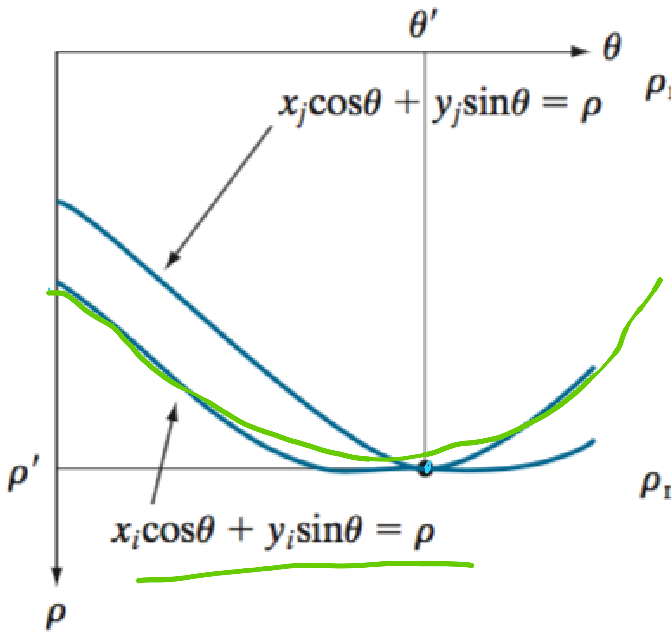

2. Create a Hough Accumulator Array (keeps the votes)

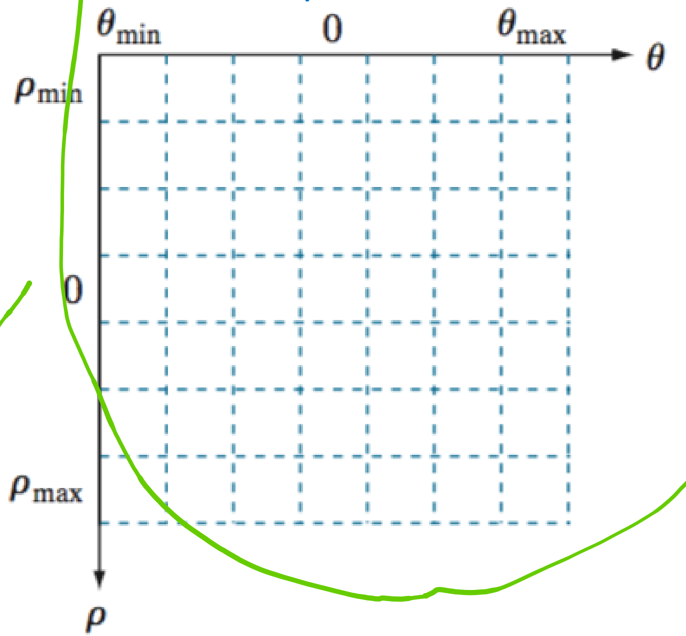

- Domain of the parametric space:

$$r \in \left[-\sqrt{M^2+N^2}, \sqrt{M^2+N^2}\right], \theta \in \left[-\frac{\pi}{2}, \frac{\pi}{2}\right]$$

M and N image resolution

Example of Hough transform

a
b

FIGURE 10.30 (a) Image of size 101×101 pixels, containing five white points (four in the corners and one in the center). (b) Corresponding parameter space.

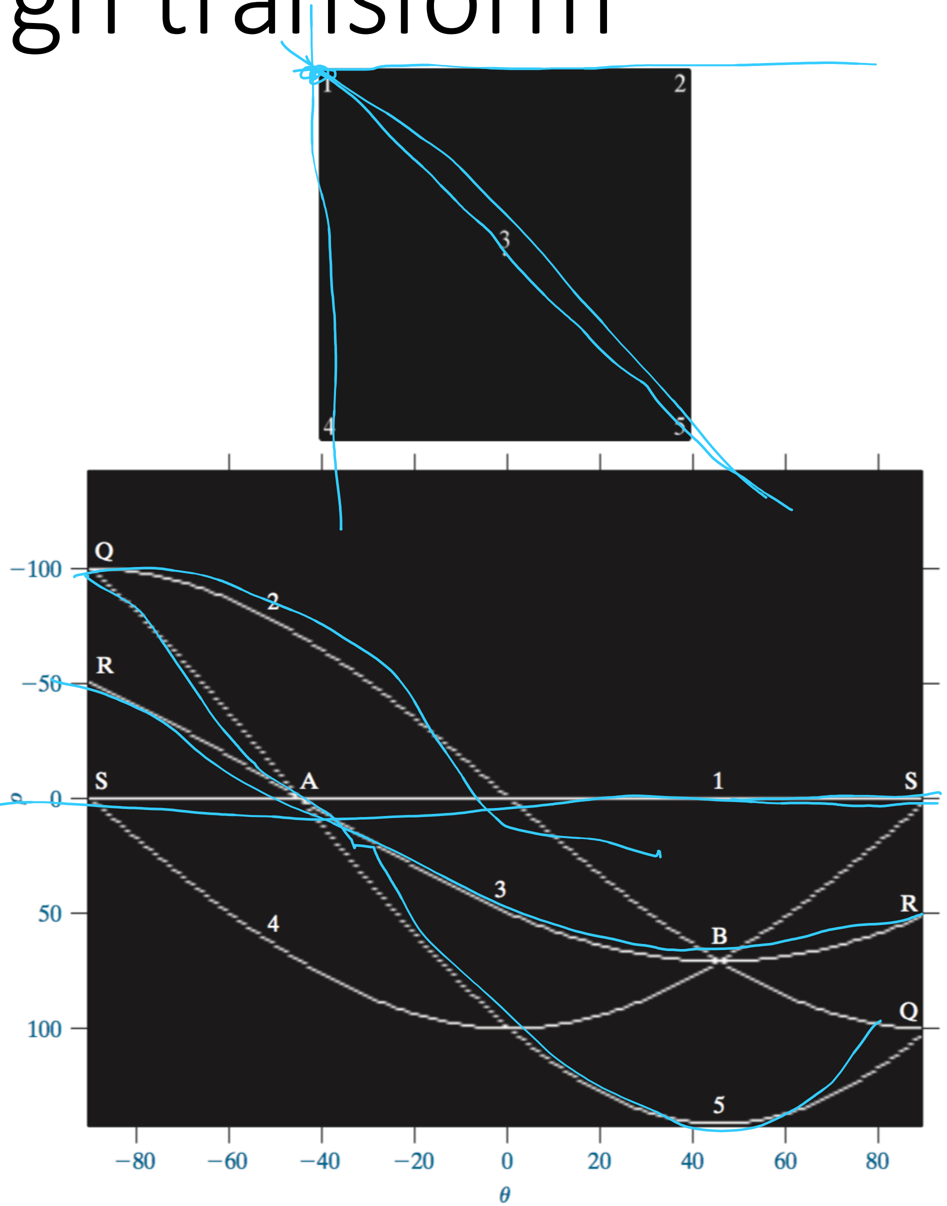

Basic Hough transform algorithm

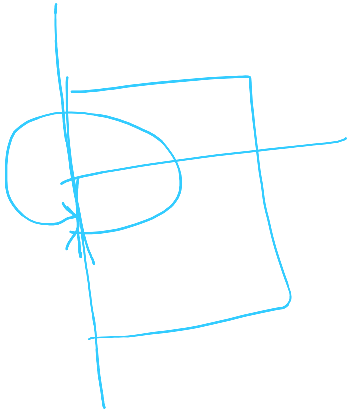

1. Initialize $H[d, \theta]=0$
2. For each **edge** point in $E(x, y)$ in the image

   for θ = 0 to 180 // some quantization; why not 2pi?

   $d = x\cos\theta + y\sin\theta$ // maybe negative

   $H[d, \theta]$ += 1
3. Find the value(s) of (d, θ) where $H[d, \theta]$ is maximum
4. The detected line in the image is given

   by $d = x\cos\theta + y\sin\theta$

Line detection example

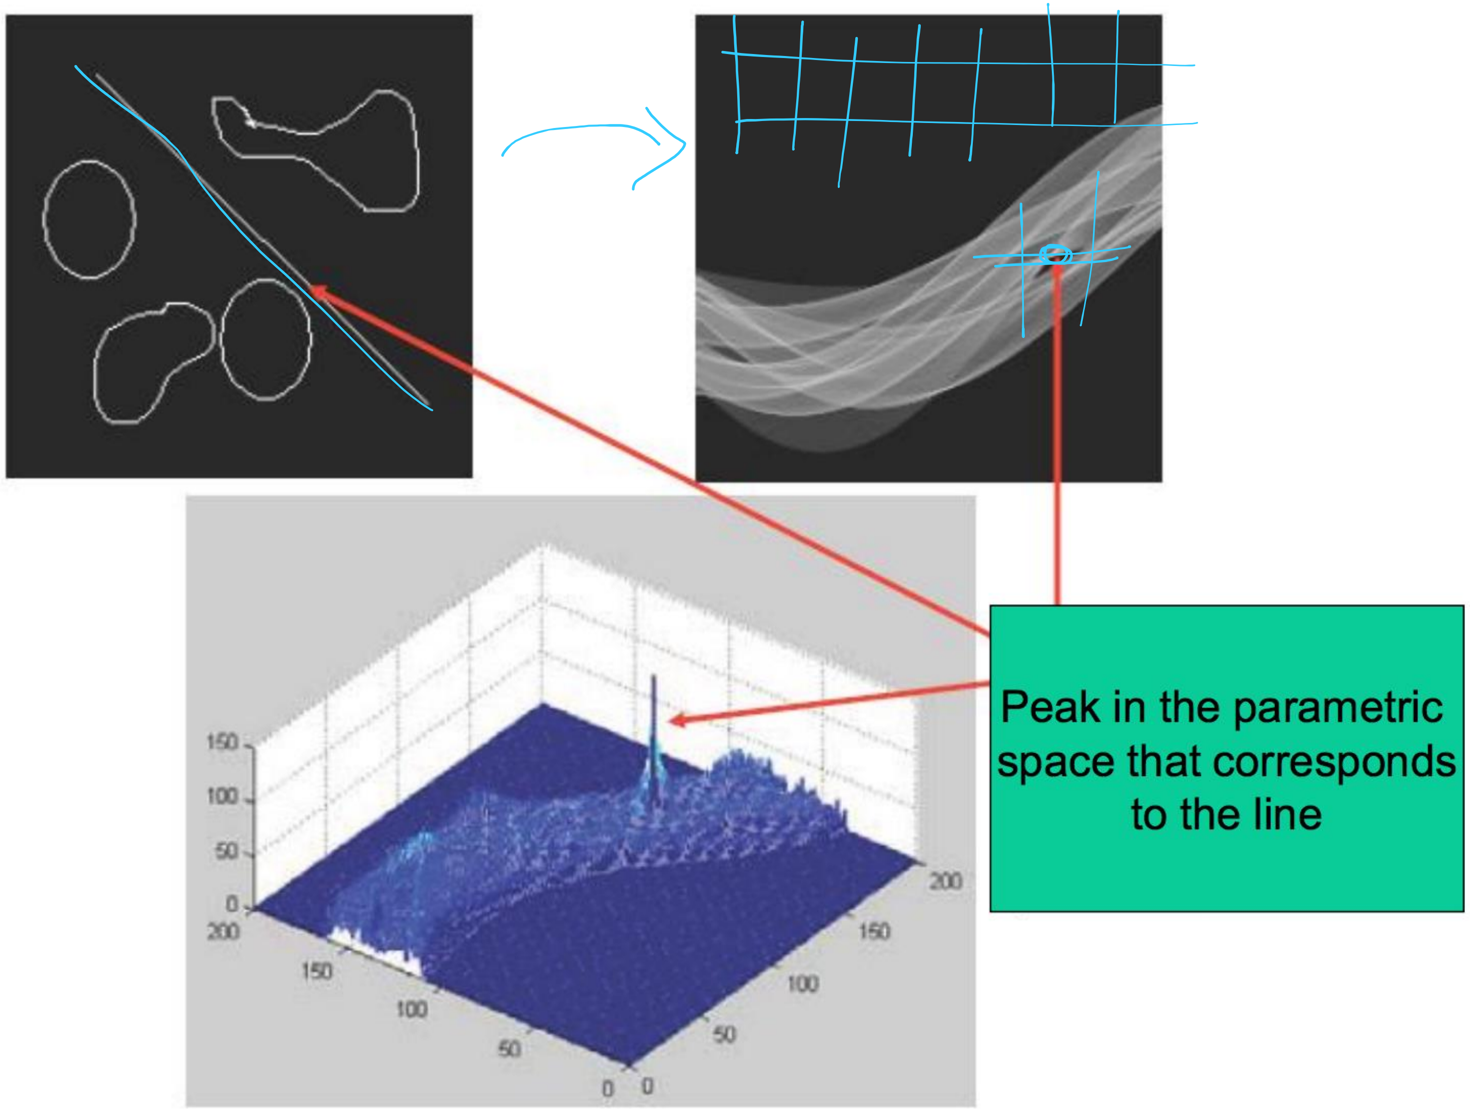

Line detection example

original

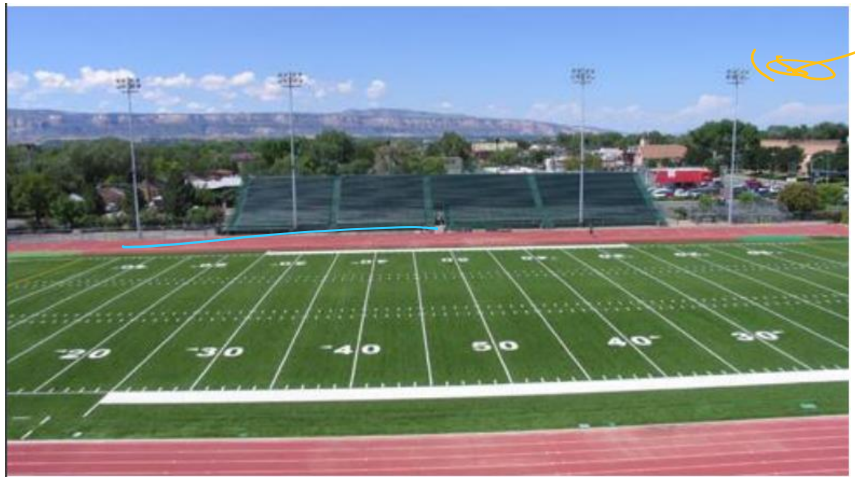

Canny edges

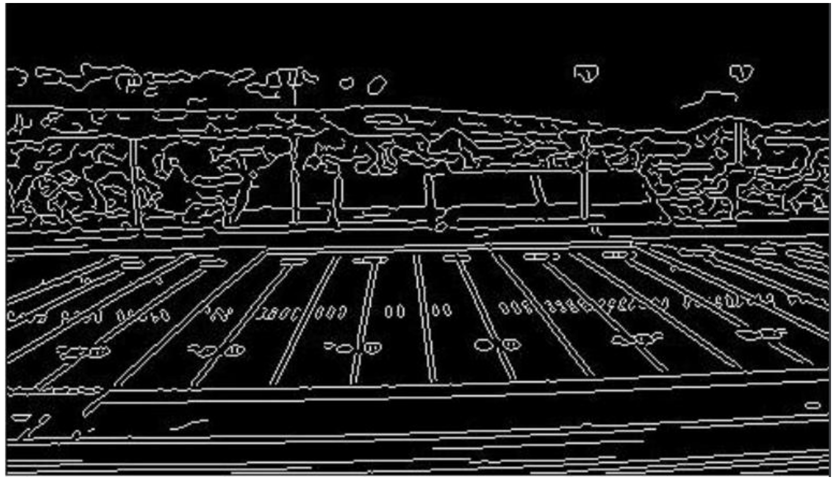

Vote space and
top peaks

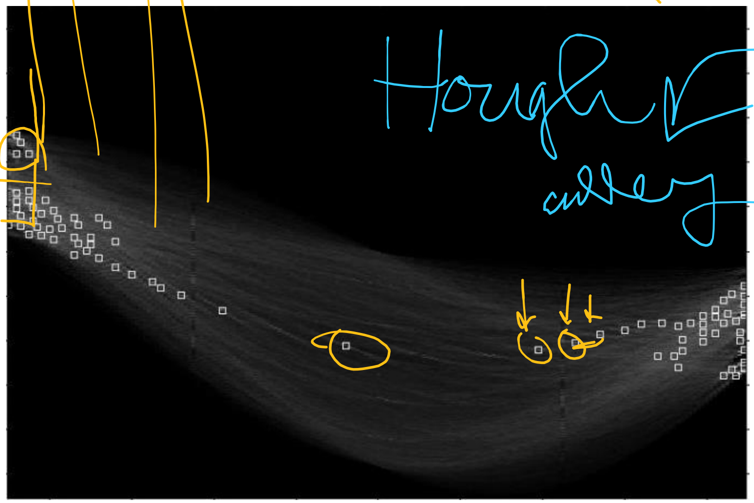

Longest segments
found

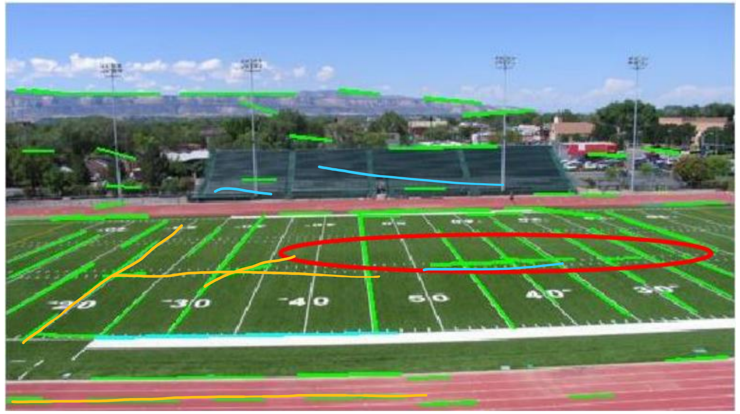

Impact of noise on Hough

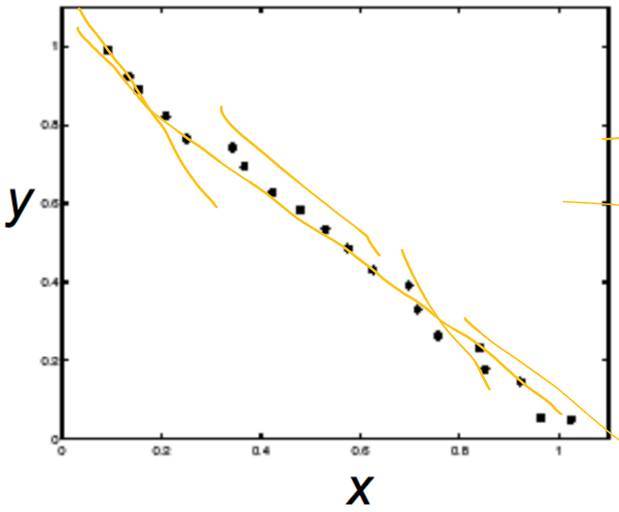

Image space
edge coordinates

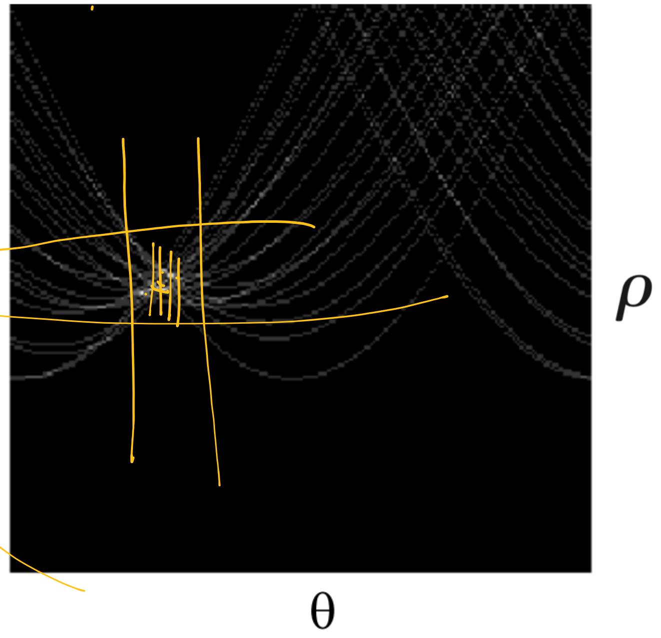

Votes

What difficulty does this present for an implementation?

Impact of noise on Hough

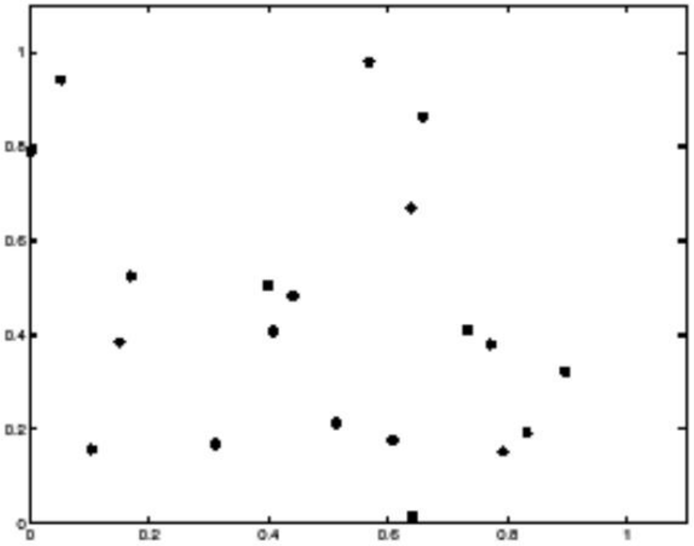

Image space edge coordinates

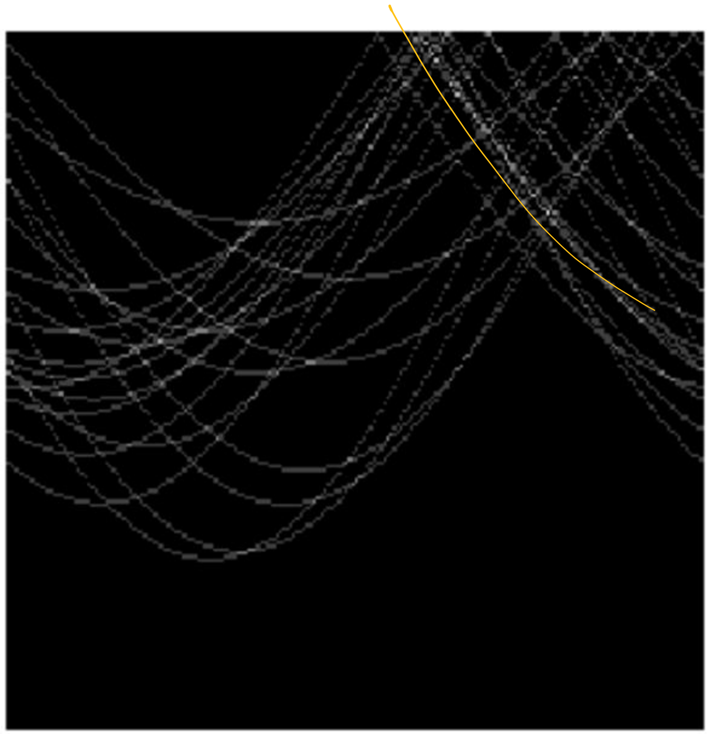

Votes

Everything appears to be “noise”, or random edge points, but we still see some peaks in the vote space

Texture analysis

What is texture?

 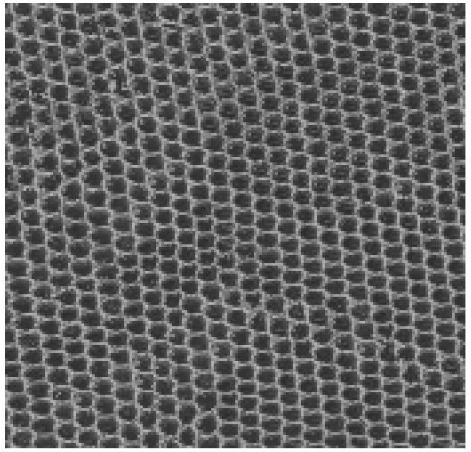 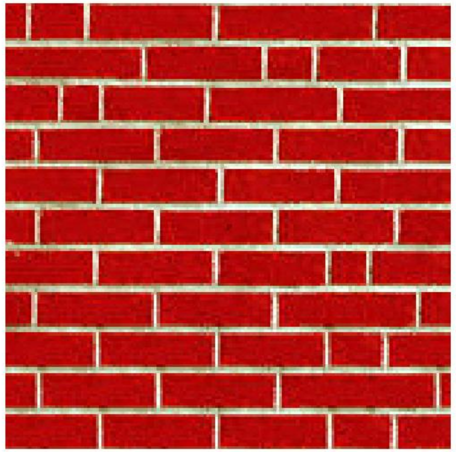

- An image obeying some statistical properties
- Similar structures repeated over and over again
- Often has some degree of randomness

Aspects of texture

- Size/granularity (sand versus pebbles versus boulders)
- Directionality/Orientation
- Random or regular (stucco versus bricks)

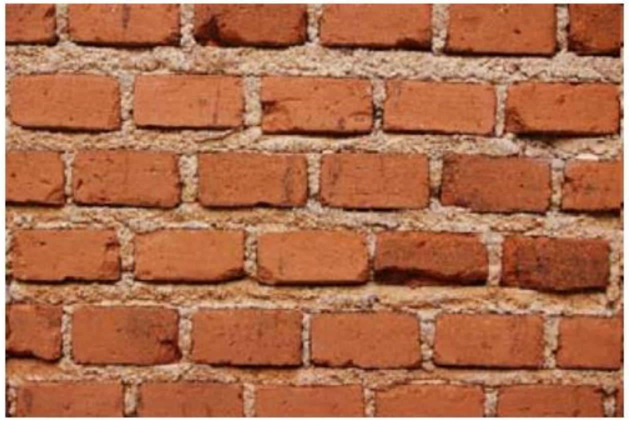

Statistical approach to texture

- Characterize texture using statistical measures computed from grayscale intensities (or colors) alone
- Less intuitive, but applicable to all images and computationally efficient
- Can be used for both classification of a given input texture and segmentation of an image into different regions

Some (simple) statistical texture measures

- Edge density and direction

- Use an edge detector as the first step in texture analysis
- The number of edge pixels in a fixed-size region tells us how busy that region is
- The directions of the edges also help characterize the texture

Two edge-based texture measures

1. edgeness per unit area

$$F_{edgeness} = |\{ p \mid \text{gradient\_magnitude}(p) \geq \text{threshold}\}| \, / \, N$$

where N is the size of the unit area

2. edge magnitude and direction histograms

$$F_{magdir} = (H_{magnitude}, H_{direction})$$

where these are the normalized histograms of gradient magnitudes and gradient directions, respectively.

Example

Original Image

Frei-Chen
Edge Image

Thresholded
Edge Image

Different F_{edgeness} for different regions

Local binary pattern measure

Multiresolution Gray Scale and Rotation Invariant Texture Classification with Local Binary Patterns

Timo Ojala, Matti Pietikäinen and Topi Mäenpää

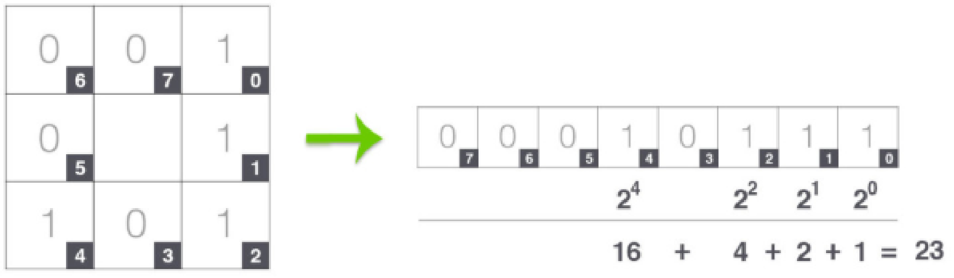

Figure 2: Taking the 8-bit binary neighborhood of the center pixel and converting it into a decimal representation. (Thanks to Bikramjot of [Hanzra Tech](#) for the inspiration on this visualization!)

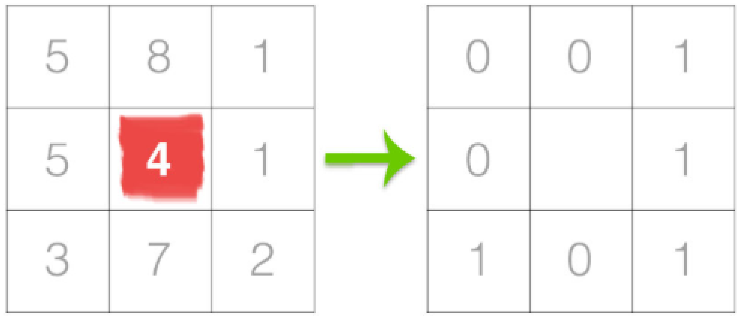

Figure 1: The first step in constructing a LBP is to take the 8 pixel neighborhood surrounding a center pixel and threshold it to construct a set of 8 binary digits.

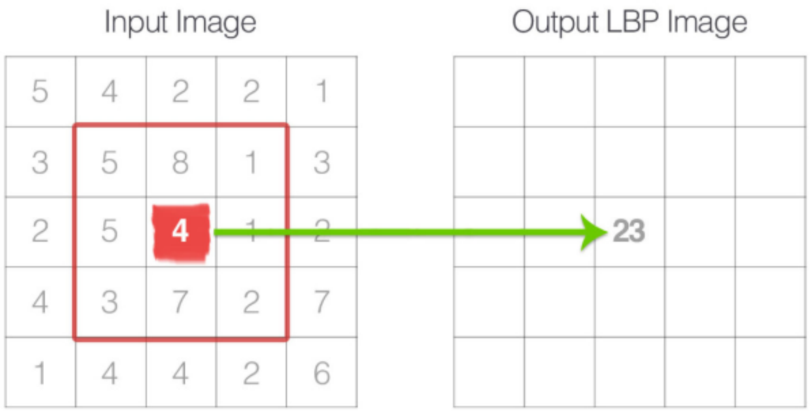

Figure 3: The calculated LBP value is then stored in an output array with the same width and height as the original image.

Image credit: https://www.pyimagesearch.com/2015/12/07/local-binary-patterns-with-python-opencv/

LBP example

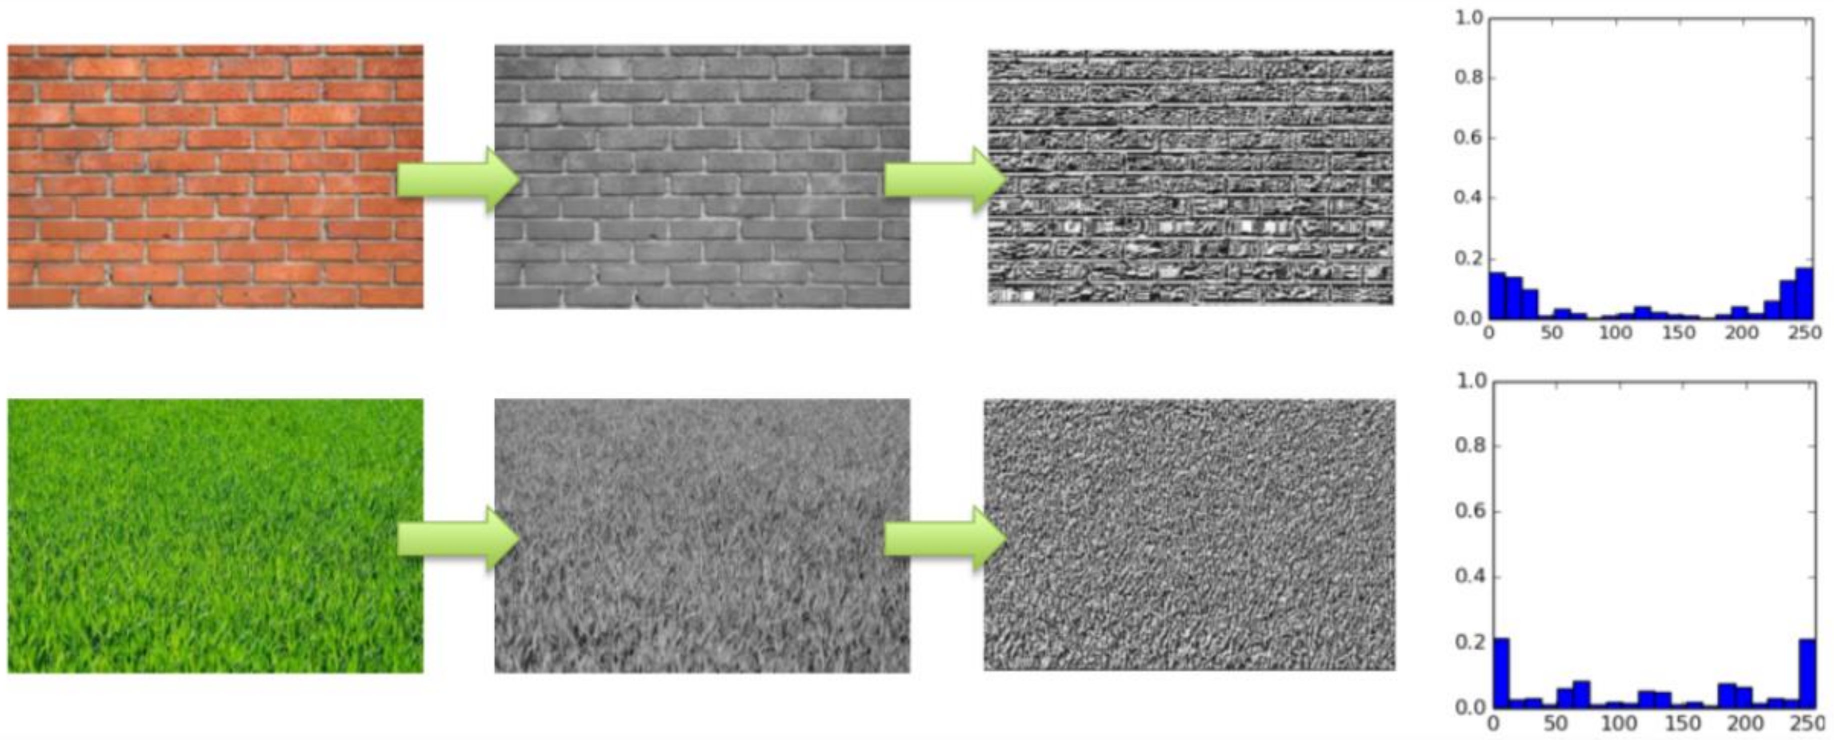

Color Image -> Grayscale Image -> LBP Mask -> Normalized LBP Histogram

Image credit: http://hanzratech.in/2015/05/30/local-binary-patterns.html

Gray Level Co-occurence Matrix (GLCM)

Textural Features for Image Classification

ROBERT M. HARALICK, K. SHANMUGAM, AND ITS'HAK DINSTEIN

- Distribution of co-occurring pixel values (grayscale values, or colors) at a given offset
  - A distance *d*, and an angle θ

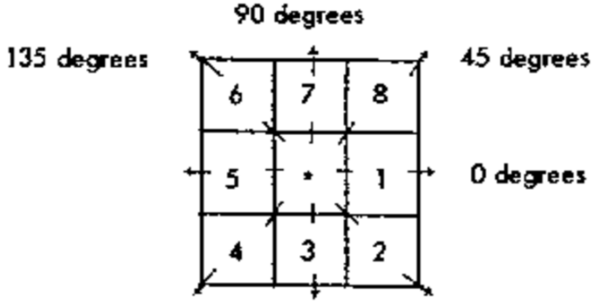

Haralick features -→ contrast, correlation texture features

GLCM example

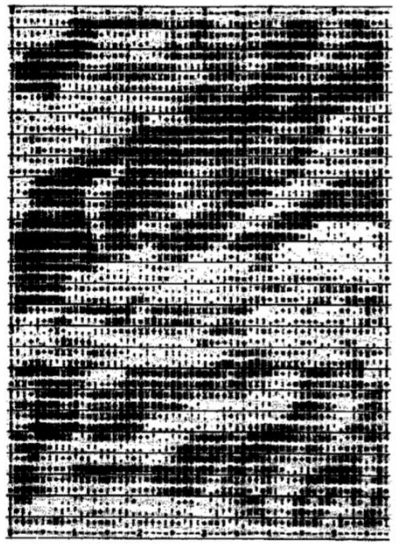

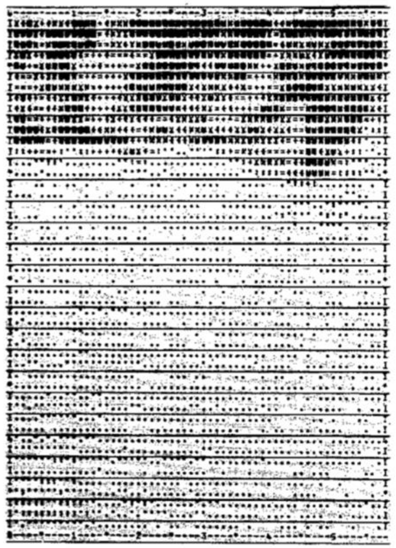

	Grassland			Water Body		
Angle	ASM	Contrast	Correlation	ASM	Contrast	Correlation
0°	.0128	3.048	.8075	.1016	2.153	.7254
45°	.0080	4.011	.6366	.0771	3.057	.4768
90°	.0077	4.014	.5987	.0762	3.113	.4646
135°	.0064	4.709	.4610	.0741	3.129	.4650
Avg.	.0087	3.945	.6259	.0822	2.863	.5327

(a) (b)

Fig. 4. Textural features for two different land-use category images.